

Importing Libraries

In [16]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

In [17]:

```
df=pd.read_csv("innovation_and_development_database.csv")
df
```

Out[17]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	NaN	NaN	
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	NaN	NaN	
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	NaN	NaN	
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	NaN	NaN	
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	NaN	NaN	
...	
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0	12
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0	12
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0	12
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	NaN	8.0	12
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	NaN	NaN	

8295 rows × 33 columns

Data Cleaning and Data Preprocessing

In [18]:

```
df=df.fillna(1)
df
```

Out[18]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0	
...	

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0	12
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0	12
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0	12
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	1.000000e+00	8.0	12
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	1.000000e+00	1.0	

8295 rows × 33 columns

In [19]:

```
df.columns
```

Out[19]: Index(['country', 'code', 'year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa', 'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro', 'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp', 'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal', 'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock', 'designpatstock', 'plantpat', 'designpat'], dtype='object')

In [20]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   country                8295 non-null   object
1   code                  8295 non-null   object
2   year                  8295 non-null   int64
3   eap                   8295 non-null   int64
4   eca                   8295 non-null   int64
5   lac                   8295 non-null   int64
6   mena                  8295 non-null   int64
7   sha                   8295 non-null   int64
8   sa                    8295 non-null   int64
9   hi                    8295 non-null   float64
10  pat                   8295 non-null   float64
11  patepo                8295 non-null   float64
12  royal                 8295 non-null   float64
13  rdexp                 8295 non-null   float64
14  rdper                 8295 non-null   float64
15  rdfinabro             8295 non-null   float64
16  rdfinprod             8295 non-null   float64
17  rdperfprod            8295 non-null   float64
18  rdperfhe              8295 non-null   float64
19  rdperfpub             8295 non-null   float64
20  lowrdexp              8295 non-null   float64
21  lowrdfinprod          8295 non-null   float64
22  lowrdperfprod         8295 non-null   float64
23  y                     8295 non-null   float64
24  stockpatEPO           8295 non-null   float64
25  poptotal              8295 non-null   float64
26  labor                 8295 non-null   float64
27  rdexpgdp              8295 non-null   float64
28  patgrantedstock       8295 non-null   float64
29  plantpatstock         8295 non-null   float64
30  designpatstock        8295 non-null   float64
31  plantpat              8295 non-null   float64
32  designpat             8295 non-null   float64
```

dtypes: float64(24), int64(7), object(2)
memory usage: 2.1+ MB

In [21]:

```
data=df[['year' , 'lac']]
data
```

Out[21]:

	year	lac
0	1960	1
1	1961	1
2	1962	1
3	1963	1
4	1964	1
...
8290	1998	0
8291	1999	0
8292	2000	0
8293	2001	0
8294	2002	0

8295 rows × 2 columns

In [25]:

```
a=df.head(60)
a
```

Out[25]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	poptotal	labor
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
5	Aruba	ABW	1965	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
6	Aruba	ABW	1966	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
7	Aruba	ABW	1967	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
8	Aruba	ABW	1968	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
9	Aruba	ABW	1969	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
10	Aruba	ABW	1970	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
11	Aruba	ABW	1971	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
12	Aruba	ABW	1972	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
13	Aruba	ABW	1973	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
14	Aruba	ABW	1974	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
15	Aruba	ABW	1975	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	poptotal	labor
16	Aruba	ABW	1976	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
17	Aruba	ABW	1977	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
18	Aruba	ABW	1978	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
19	Aruba	ABW	1979	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
20	Aruba	ABW	1980	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
21	Aruba	ABW	1981	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
22	Aruba	ABW	1982	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
23	Aruba	ABW	1983	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
24	Aruba	ABW	1984	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
25	Aruba	ABW	1985	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
26	Aruba	ABW	1986	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
27	Aruba	ABW	1987	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
28	Aruba	ABW	1988	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
29	Aruba	ABW	1989	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
30	Aruba	ABW	1990	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
31	Aruba	ABW	1991	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
32	Aruba	ABW	1992	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
33	Aruba	ABW	1993	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
34	Aruba	ABW	1994	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
35	Aruba	ABW	1995	0	0	1	0	0	0	0.0	...	1.0	0.0	1.0	1.0
36	Aruba	ABW	1996	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
37	Aruba	ABW	1997	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
38	Aruba	ABW	1998	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
39	Aruba	ABW	1999	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
40	Aruba	ABW	2000	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
41	Aruba	ABW	2002	0	0	1	0	0	0	0.0	...	1.0	1.0	1.0	1.0
42	Andorra	ADO	1960	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
43	Andorra	ADO	1961	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
44	Andorra	ADO	1962	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
45	Andorra	ADO	1963	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
46	Andorra	ADO	1964	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
47	Andorra	ADO	1965	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
48	Andorra	ADO	1966	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
49	Andorra	ADO	1967	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
50	Andorra	ADO	1968	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
51	Andorra	ADO	1969	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO	poptotal	labor
52	Andorra	ADO	1970	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
53	Andorra	ADO	1971	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
54	Andorra	ADO	1972	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
55	Andorra	ADO	1973	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
56	Andorra	ADO	1974	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
57	Andorra	ADO	1975	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
58	Andorra	ADO	1976	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0
59	Andorra	ADO	1977	0	1	0	0	0	0	0.0	...	1.0	1.0	1.0	1.0


60 rows × 33 columns

Line chart

In [26]:

data.plot.line(subplots=True)

Out[26]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)




Line chart

In [27]:

data.plot.line()

Out[27]: <AxesSubplot:>




Bar chart

```
In [28]: b=data[0:50]
```

```
In [29]: b.plot.bar()
```


```
Out[29]: <AxesSubplot:>
```



Histogram

```
In [30]: data.plot.hist()
```

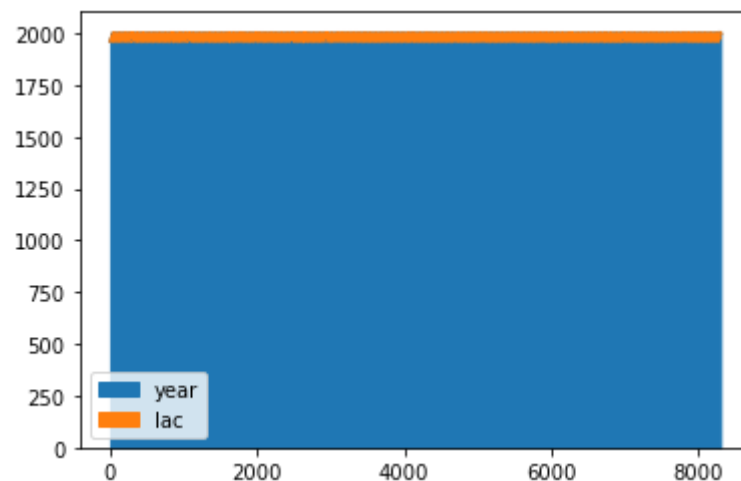
```
Out[30]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [31]: `data.plot.area()`


Out[31]: <AxesSubplot:>



Box chart

In [32]: `data.plot.box()`

Out[32]: <AxesSubplot:>



Pie chart

```
In [33]: b.plot.pie(y='lac' )
```


```
Out[33]: <AxesSubplot:ylabel='lac'>
```




Scatter chart

```
In [34]: data.plot.scatter(x='year' ,y='lac')
```

Out[34]: <AxesSubplot:xlabel='year', ylabel='lac'>



In [35]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   country              8295 non-null   object
1   code                 8295 non-null   object
2   year                 8295 non-null   int64
3   eap                  8295 non-null   int64
4   eca                  8295 non-null   int64
5   lac                  8295 non-null   int64
6   mena                 8295 non-null   int64
7   sha                  8295 non-null   int64
8   sa                   8295 non-null   int64
9   hi                   8295 non-null   float64
10  pat                  8295 non-null   float64
11  patepo               8295 non-null   float64
12  royal                8295 non-null   float64
13  rdexp                8295 non-null   float64
14  rdper                8295 non-null   float64
15  rdfinabro            8295 non-null   float64
16  rdfinprod            8295 non-null   float64
17  rdperfprod           8295 non-null   float64
18  rdperfhe             8295 non-null   float64
19  rdperfpub            8295 non-null   float64
20  lowrdexp             8295 non-null   float64
21  lowrdfinprod         8295 non-null   float64
22  lowrdperfprod        8295 non-null   float64
23  y                    8295 non-null   float64
24  stockpatEPO          8295 non-null   float64
25  poptotal             8295 non-null   float64
26  labor                8295 non-null   float64
27  rdexpgdp             8295 non-null   float64
28  patgrantedstock      8295 non-null   float64
29  plantpatstock        8295 non-null   float64
30  designpatstock       8295 non-null   float64
31  plantpat             8295 non-null   float64
32  designpat            8295 non-null   float64
dtypes: float64(24), int64(7), object(2)
memory usage: 2.1+ MB
```

In [38]:

df.columns

```
Out[38]: Index(['country', 'code', 'year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',
                'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',
                'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',
```

```
'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEP0', 'poptotal',  
'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',  
'designpatstock', 'plantpat', 'designpat'],  
dtype='object')
```

In [36]:

df.describe()

Out[36]:

	year	eap	eca	lac	mena	sha	sa	
count	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8
mean	1981.203014	0.094515	0.195901	0.176974	0.098614	0.150090	0.026522	
std	12.421590	0.292561	0.396917	0.381670	0.298161	0.357182	0.160691	
min	1960.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1970.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	1981.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	1992.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
max	2002.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

8 rows × 31 columns



In [39]:


df1=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa','hi']]

EDA AND VISUALIZATION

In [40]:


sns.pairplot(df1[0:50])

Out[40]: <seaborn.axisgrid.PairGrid at 0x2a1696d65b0>




In [44]: `sns.displot(df1['year'])`

Out[44]: `<seaborn.axisgrid.FacetGrid at 0x2a16f3d6a90>`



```
In [42]: sns.heatmap(df1.corr())
```

```
Out[42]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [45]: x=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa']]
          y=df['hi']
```

```
In [46]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [47]: from sklearn.linear_model import LinearRegression
          lr=LinearRegression()
```

```
lr.fit(x_train,y_train)
```

Out[47]: LinearRegression()

```
In [48]: lr.intercept_
```

Out[48]: -0.10129916868360311

```
In [49]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```


Out[49]:

	Co-efficient
--	--------------

year	0.000312
eap	-0.516784
eca	-0.516617
lac	-0.516767
mena	-0.516982
sha	-0.516911
sa	-0.516918

```
In [50]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[50]: <matplotlib.collections.PathCollection at 0x2a16fc09d00>



ACCURACY

```
In [51]: lr.score(x_test,y_test)
```

Out[51]: 0.4613696470505868

```
In [52]: lr.score(x_train,y_train)
```

Out[52]: 0.442752791624356

Ridge and Lasso

```
In [53]: from sklearn.linear_model import Ridge,Lasso
```

```
In [54]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[54]: Ridge(alpha=10)

Accuracy(Ridge)

```
In [55]: rr.score(x_test,y_test)
```

Out[55]: 0.458186877260473

```
In [56]: rr.score(x_train,y_train)
```

Out[56]: 0.44143721604040387

```
In [57]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[57]: Lasso(alpha=10)

```
In [58]: la.score(x_train,y_train)
```

Out[58]: 0.0

Accuracy(Lasso)

```
In [59]: la.score(x_test,y_test)
```

Out[59]: -0.00034648662632985605

```
In [60]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[60]: ElasticNet()

```
In [61]: en.coef_
```

Out[61]: array([0., -0., -0., -0., -0., -0., -0.])

```
In [62]: en.intercept_
```

```
Out[62]: 0.13296589734757147
```

```
In [63]: prediction=en.predict(x_test)
```

```
In [64]: en.score(x_test,y_test)
```

```
Out[64]: -0.00034648662632985605
```

Evaluation Metrics

```
In [65]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
0.23530485566046233
0.12001888817033582
0.3464374231666317
```

Logistic Regression

```
In [66]: from sklearn.linear_model import LogisticRegression
```

```
In [67]: feature_matrix=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa']]
target_vector=df['hi']
```

```
In [68]: feature_matrix.shape
```

```
Out[68]: (8295, 7)
```

```
In [69]: target_vector.shape
```

```
Out[69]: (8295,)
```

```
In [70]: from sklearn.preprocessing import StandardScaler
```

```
In [71]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [72]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[72]: LogisticRegression(max_iter=10000)
```



```
In [74]: observation=[[1,2,3,4,5,6,7]]
```

```
In [75]: prediction=logr.predict(observation)
print(prediction)

[0.]
```

```
In [76]: logr.classes_
```

```
Out[76]: array([0., 1.])
```

```
In [77]: logr.score(fs,target_vector)
```

```
Out[77]: 0.8772754671488848
```

```
In [78]: logr.predict_proba(observation)[0][0]
```

```
Out[78]: 1.0
```

```
In [79]: logr.predict_proba(observation)
```

```
Out[79]: array([[1.00000000e+00, 8.32677247e-28]])
```

Random Forest

```
In [80]: from sklearn.ensemble import RandomForestClassifier
```

```
In [81]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[81]: RandomForestClassifier()
```

```
In [82]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]
                }
```

```
In [83]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[83]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

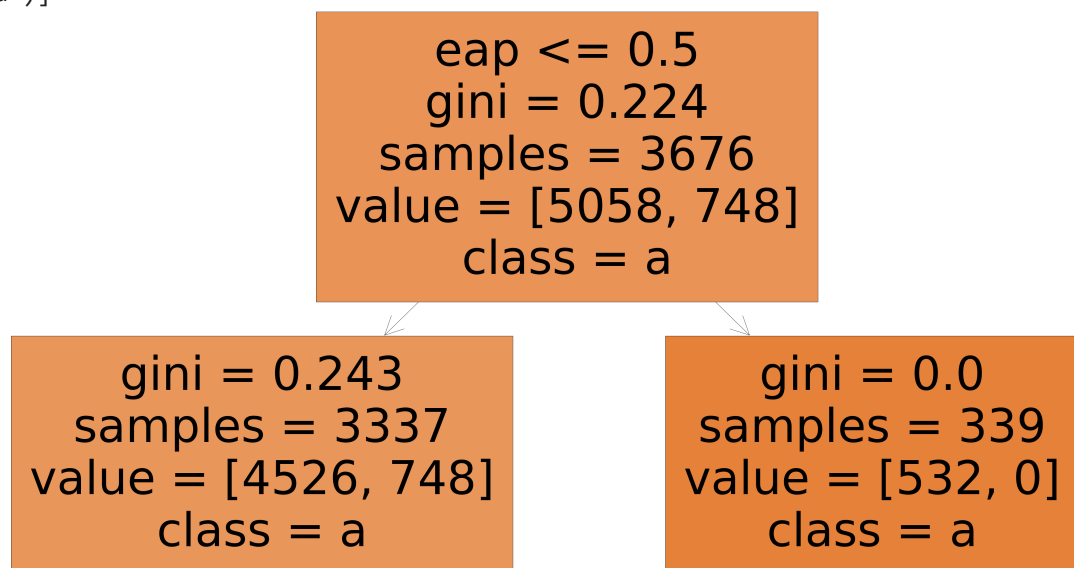
```
In [84]: grid_search.best_score_
```

Out[84]: 0.8670341026524285

In [85]: `rfc_best=grid_search.best_estimator_`

In [86]: `from sklearn.tree import plot_tree`
`plt.figure(figsize=(80,40))`
`plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c'],`

Out[86]: [Text(2232.0, 1630.8000000000002, 'eap <= 0.5\ngini = 0.224\nsamples = 3676\nvalue = [5058, 748]\nnclass = a'),
 Text(1116.0, 543.5999999999999, 'gini = 0.243\nsamples = 3337\nvalue = [4526, 748]\nnclass = a'),
 Text(3348.0, 543.5999999999999, 'gini = 0.0\nsamples = 339\nvalue = [532, 0]\nnclass = a')]



Conclusion

Accuracy

Linear Regression

In [87]: `lr.score(x_test,y_test)`

Out[87]: 0.4613696470505868

In [88]: `lr.score(x_train,y_train)`

Out[88]: 0.442752791624356

Lasso

In [89]: `la.score(x_test,y_test)`

Out[89]: -0.00034648662632985605

In [90]: `la.score(x_train,y_train)`

Out[90]: 0.0

Ridge

In [91]: `rr.score(x_test,y_test)`

Out[91]: 0.458186877260473

In [92]: `rr.score(x_train,y_train)`

Out[92]: 0.44143721604040387

Elastic Net

In [93]: `en.score(x_test,y_test)`

Out[93]: -0.00034648662632985605

Logistic Regression

In [95]: `logr.score(fs,target_vector)`

Out[95]: 0.8772754671488848

Random Forest

In [96]: `grid_search.best_score_`

Out[96]: 0.8670341026524285

From the above data, we can conclude that logistic regression is preferable to other regression types

In []:

In []: